

Can we boot Linux from just a floppy?

An exercise in minimalism.

Kamila Szewczyk, *Association for Computational Heresy*

19 Sesja Linuksowa, Wrocław, April 2025

Introduction

- Some common storage media.



(a) CD-ROM



(b) USB drive



(c) SD Card

- Their capacities are usually overwhelming and measured in gigabytes (10^9 bytes).
- All of them constitute popular boot and installation media for Linux.
- Do we really need this much storage just to boot an operating system?

Introduction

- Prior art:
 - Damn Small Linux (DSL) - 50MB - A live CD distribution for 486+ CPUs, 8MB of RAM.
 - Tiny Core Linux - 17MB - Perhaps the most popular minimalist distribution.
 - NanoLinux - 14MB - Discontinued project.
- All of those contenders are quite small and still provide many interesting features like web browsers, text editors or even games.
- They also lag behind in kernel versions.
- What if we didn't want *any* of the nice features, and simply wanted a minimal TTY environment that still formally runs Linux?

Introduction

- As a goalpost of our experiment, we will use this magical device:



Figure: A 3.5" floppy disk.

- Each of these holds about 1.44MB of data, about 0.2% of a standard 700MB CD-ROM.
- The best widespread contender, Tiny Core Linux, needs at least 10 of these to fit.

- It is clear that we will have to compile the kernel and the userland ourselves.
- Because 32-bit x86 code tends to be smaller than its 64-bit counterpart, we will build a i486+ kernel.
 - This has some more advantages: we turn off fragments of the kernel code responsible for complicated hardware instructions that the 486 simply doesn't have.
- But we don't want just *any* 486 ELF toolchain: glibc is much too bulky for our needs.
- Instead, we will use musl, a lightweight C library, which tends to be much smaller and compatible with glibc on many counts.
- Equipped with i486-linux-musl-gcc, we may continue.
- In case you want to follow along, you can get this from <https://musl.cc/>.

Toolchain: experiment

```
% cat hello.c
#include <stdio.h>
int main(void) puts("Hello, world!"); for (;;) ;
% i486-linux-musl-gcc hello.c -o hello -Os -static -s
% wc -c hello && ./hello
4948 hello
Hello, world!
% ldd hello
not a dynamic executable
```

- The resulting 32-bit program only uses system calls and does not depend on any shared libraries.
- 5KB for a simple program is not fantastic, but it is a good start.
- Let's not get off track. We will get back to that later.

Kernel & initrd

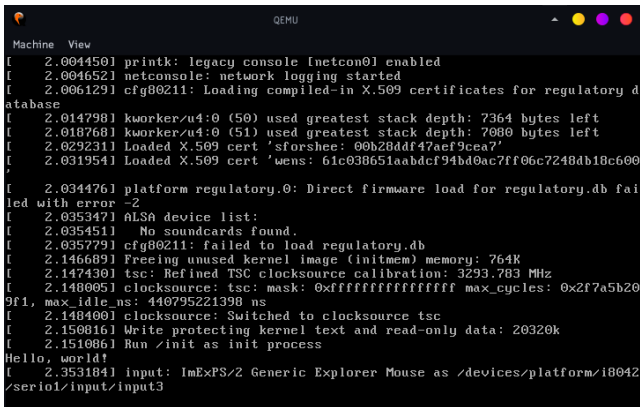
- Let's try to get something booting first.
- The specific kernel that we will be working with is `linux-6.13.8` - most recent version at the time of writing.
- Compiling the kernel by itself is pretty basic nitty-gritty.
 - Use `arch/x86/configs/i386_defconfig` as a base.
 - Enable TTY and `printk` support (for the latter you need expert settings enabled).
 - Then:

```
make ARCH=x86 CROSS_COMPILE=i486-linux-musl- bzImage.
```
- Before we can compile and boot the kernel, we need to prepare the `initrd`.
- The *Hello, world!* program is a good initial candidate: together with a small bootloader, it will leave us with a lot of space for the kernel.
- We put the `init` program in a CPIO archive, which is a standard format for `initrd` images.

Kernel & initrd

- Boot with the following command:

```
qemu-system-i386 -kernel arch/x86/boot/bzImage \
                 -initrd ./rootfs.cpio.gz
```



```
Machine View
[ 2.004450] printk: legacy console [netcon0] enabled
[ 2.004652] netconsole: network logging started
[ 2.006129] cfg80211: Loading compiled-in X.509 certificates for regulatory database
[ 2.014798] kworker/u4:0 (50) used greatest stack depth: 7364 bytes left
[ 2.018768] kworker/u4:0 (51) used greatest stack depth: 7080 bytes left
[ 2.029231] Loaded X.509 cert 'sforshee: 00b28ddf47aef9cea7'
[ 2.031954] Loaded X.509 cert 'wens: 61c038651aabdcf94bd0ac7ff06c7248db18c600'
[ 2.034476] platform regulatory.0: Direct firmware load for regulatory.db failed with error -2
[ 2.035347] ALSA device list:
[ 2.035451]   No soundcards found.
[ 2.035779] cfg80211: failed to load regulatory.db
[ 2.146689] Freeing unused kernel image (initmem) memory: 764K
[ 2.147430] tsc: Refined TSC clocksource calibration: 3293.783 MHz
[ 2.148005] clocksource: tsc: mask: 0xffffffffffffffff max_cycles: 0x2f7a5b209f1, max_idle_ns: 440795221398 ns
[ 2.148400] clocksource: Switched to clocksource tsc
[ 2.150816] Write protecting kernel text and read-only data: 20320k
[ 2.151086] Run /init as init process
Hello, world!
[ 2.353184] input: ImExPS/2 Generic Explorer Mouse as /devices/platform/i8042/serio1/input/input3
```


Kernel & initrd

- Excellent! We have a working kernel and initrd.
- Bad news: the kernel image is 10.8 MiB.
- Before we start trying to remove stuff from the kernel, let's try to build a functional initrd using BusyBox. Then it's easier to tell if the stuff we are removing is actually used by something.
- Specifically, we will use BusyBox 1.35.0. The binary is quite large, at around 1.0 MiB, but it is statically linked and contains a lot of useful tools.

Kernel & initrd

- init program: install BusyBox, mount devtmpfs and run /bin/init.

```
int main(int argc, char *argv[]) {
    pid_t pid = fork();
    if (pid == 0) {
        char * args[] = { "/bin/busybox", "--install", "-s", NULL };
        execv(args[0], args);
        perror("execv failed"), exit(1);
    } else if (pid < 0)
        perror("fork failed"), exit(1);
    int status; waitpid(pid, &status, 0);
    if (!WIFEXITED(status))
        perror("waitpid"), exit(1);
    if (mount("none", "/dev", "devtmpfs", 0, "")) perror("mount"), exit(1);
    if (mount("none", "/proc", "proc", 0, "")) perror("mount"), exit(1);
    if (mount("none", "/sys", "sysfs", 0, "")) perror("mount"), exit(1);
    // /bin/init is dmesg -n 1;exec sh
    char * args[] = { "/bin/busybox", "sh", "/bin/init", NULL };
    execv(args[0], args); perror("execv"); exit(1);
}
```

Kernel & initrd

```
QEMU
Machine View
[ 1.899459] cfg80211: failed to load regulatory.db
[ 1.900149] No soundcards found.
[ 2.021811] Freeing unused kernel image (initmem) memory: 764K
[ 2.023752] Write protecting kernel text and read-only data: 20320k
[ 2.023965] Run /init as init process
[ 2.109594] busybox (52) used greatest stack depth: 6400 bytes left
sh: can't access tty; job control turned off
/ # ls -la
total 20
drwxrwxr-x  9 1000    1000          240 Mar 23 21:30 .
drwxrwxr-x  9 1000    1000          240 Mar 23 21:30 ..
-rw-----  1 0        0              7 Mar 23 21:30 .ash_history
drwxrwxr-x  2 1000    1000        1960 Mar 23 21:30 bin
drwxr-xr-x  7 0        0        2260 Mar 23 21:30 dev
-rwxrwxr-x  1 1000    1000    13144 Mar 23 21:26 init
lrwxrwxrwx  1 0        0           12 Mar 23 21:30 linuxrc -> /bin/busybox
dr-xr-xr-x 105 0        0           0 Mar 23 21:30 proc
drwx-----  2 0        0           40 Mar 23 19:38 root
drwxrwxr-x  2 1000    1000        1480 Mar 23 21:30 sbin
dr-xr-xr-x 12 0        0           0 Mar 23 21:30 sys
drwxrwxr-x  4 1000    1000         80 Mar 23 21:04 usr
/ # uname -a
Linux (none) 6.13.8 #3 SMP PREEMPT_DYNAMIC Sun Mar 23 20:43:41 CET 2025 i686 GNU
/Linux
/ # _
```

Stripping the kernel

- We have a Linux system that fits in about 11 MiB. The CPIO with the `initrd` is about 699 KiB: we will have to look into that, because surely we can improve on it.
- The next two steps are compressing these two components well enough to fit in a total of 1.44 MB. We will start with the kernel, as it's the biggest offender.
 - First step: Use XZ compression, disable the support for all other codecs except `gzip` for `initrd`, compile with `-Os`.
 - These changes alone bring the kernel down to about 6.7 MiB.
- Unfortunately, together with the `initrd` we are still quite far away from our goal:

```
6763008 arch/x86/boot/bzImage
669206 rootfs.cpio.gz
7432214 total
```

Stripping the kernel

- We will squeeze the kernel as much as possible by removing device drivers, filesystems, networking features, debugging options, and other features that we don't need.
- Further, the build script supports embedding the initrd in the kernel image. We will use xz compression (keeping in mind that we need `--check=crc32` for the kernel to boot).
- The result? 1282560 arch/x86/boot/bzImage - the kernel + initrd fit in 1.28MB, much below the 1.44MB limit.

Stripping the kernel

- The result? 1282560 arch/x86/boot/bzImage - the kernel + initrd fit in 1.28MB, much below the 1.44MB limit.
- That's very nice, but we can't really do anything with our operating system other than idling in the shell.
- I personally don't like vi and would prefer a different text editor. Also, I would like to put some cool programs in the initramfs to play with - that would be best achieved with a dynamically linked libc.
- Also, we haven't actually figured out how to make this boot off an actual floppy (yet).

- Now that we have a MVP, we want to accomplish the following:
 - ① Build a bootable floppy disk instead of relying on `qemu-system-i386 -kernel arch/x86/boot/bzImage` to boot.
 - ② Strip down busybox to remove things that we won't need. For example user management, various `mkfs/fsck` programs for filesystems that we don't support.
 - ③ Add more programs to the `initramfs` and features to the kernel: interpreters, compilers, games, etc.
- Of course, eventually we will run out of space. Then we will look into custom compression.

Adding networking

- We will enable networking, E1000 support, the TCP/IP stack and the basic security features.
- This unfortunately makes the kernel swell up quite a bit.
- Configuring networking with busybox is tricky. I wrote the following script:

```
#!/bin/sh
ip link set eth0 up
LEASE=$(udhcpc -i eth0 -n 2>&1)
IP=$(echo "$LEASE" | awk '/lease of/ {print $4}')
GATEWAY=$(echo "$LEASE" | awk '/obtained from/ {print $7}' | tr -d ',')
if [ -n "$IP" ] && [ -n "$GATEWAY" ]; then
    ip addr flush dev eth0
    ip addr add "$IP/24" dev eth0
    ip route add default via "$GATEWAY"
    echo "Network configured: IP=$IP, Gateway=$GATEWAY"
else
    echo "Failed to obtain an IP address via DHCP." && exit 1
fi
```


Networking

```
QEMU
Machine View
e1000: Intel(R) PRO/1000 Network Driver
e1000: Copyright (c) 1999-2006 Intel Corporation.
e1000 0000:00:03.0: found PCI INT A -> IRQ 11
e1000 0000:00:03.0 eth0: (PCI:33MHz:32-bit) 52:54:00:12:34:56
e1000 0000:00:03.0 eth0: Intel(R) PRO/1000 Network Connection
serio: i8042 KBD port at 0x60,0x64 irq 1
serio: i8042 AUX port at 0x60,0x64 irq 12
NET: Registered PF_INET6 protocol family
input: AT Translated Set 2 keyboard as /devices/platform/i8042/serio0/input/input0
Segment Routing with IPv6
In-situ OAM (IOAM) with IPv6
sit: IPv6, IPv4 and MPLS over IPv4 tunneling driver
NET: Registered PF_PACKET protocol family
sched_clock: Marking stable (474464875, 9519122)->(490104860, -6120863)
Freeing unused kernel image (initmem) memory: 1276K
Write protecting kernel text and read-only data: 3004k
Run /init as init process
sh: can't access tty: job control turned off
/ # sh setup-net.sh
Network configured: IP=10.0.2.15, Gateway=10.0.2.2
/ # ping 1.1.1.1
PING 1.1.1.1 (1.1.1.1): 56 data bytes
64 bytes from 1.1.1.1: seq=0 ttl=255 time=51.267 ms
```

Checkpoint: list of features

- `printk` for debugging, ISA bus support, R/W block layer, compacting memory manager, PCI bus support.
- SATA/ATA, SCSI, E1000 (generic network card interface) support, IPv4 networking, DNS resolution, FAT32 file systems.
- BusyBox with a DHCP network setup script, `dmesg`, `wget`, `vi`, `sh`, `awk`, `dc`, `bc`, `httpd`, and so on.
- Caveat: the XZ image is 1'835'520. We will implement a custom compression algorithm to fit this in a floppy.

Kernel Compression

- Basic idea: we put the initrd and kernel in the same logical container, which is then bootstrapped by `mkpiggy` (a custom program that Linux uses for compressed kernels).
- Nominally, this is simple on paper: we start with some existing compression wrapper (e.g. XZ), blank out the source code, change the makefile to use a different compressor, then embed the decompressor in the kernel.
- Issues:
 - We are looking for a somewhat fast statistical compressor that can beat XZ quite significantly.
 - We would prefer it to not use gobs of memory and to be able to run on a 486.
 - It needs to fit in the piggyback module together with the compressed kernel and bootloader, i.e. it can not be too large.

Kernel Compression



The work is mysterious and important.

Kernel Compression

- We will reuse some common ideas: PAQ-like context-adaptive binary arithmetic coding, a context model with a rudimentary context mixer between finite order context models and a match model that seldom disables it.
- The mixer output is filtered by a chain of a few APMs. The input is filtered by an E8E9 transform together with a simple disassembly pass that recognises prefixes and opcodes.
- The biggest issue is integrating this with the Linux kernel, which is somewhat unfamiliar to me as a C code base.
- If you want to know what any of this means, it will probably be explained in my compression book that I talked about here about a year ago.

Kernel Compression

- To not go into too much detail, the decompression stub was easily produced in a few hours.
- It's somewhat fast (20s to decode the kernel) on my machine.
- Most importantly, the kernel now fits into 1.44MB!

1413632 arch/x86/boot/bzImage

- Two more goal posts to score: a standalone boot loader (\Rightarrow a working bootable floppy image) and some cool stuff in the user land if we can squeeze it in.

The Bootloader

- Written from scratch in x86 assembly, about 300 lines of code.
- Core difficulties are:
 - Enabling the A20 gate (requires some trickery because of how many disjoint methods there are - via BIOS, via keyboard, etc).
 - Operating in Unreal Mode so we can use BIOS interrupts to move stuff from conventional memory containing sectors read to a high memory address where the kernel expects to be loaded.
 - Loading the GDT and properly invoking the kernel, but that's easy.
- It ends up taking less than two sectors in total even with my sloppy and rushed programming.

End Result

- Direct download (1.44M floppy image):
 - <https://palaiologos.rocks/doc/projects/floppy-linux.img>
- This also version features `fasm` (Flat Assembler by Tomek Grysztar) in the `initrd`.
- Adding an alternative program that is about 120KB in size was also possible, but it quite like `fasm` and don't see the point in doing that.

End Result: Full list of features (I)

- All of BusyBox; abbreviated list below:
 - Shell utilities: the ash shell, the hush shell, base32, base64, chattr, chmod, cp, date, dd, df, egrep, fgrep, getopt, iostat, kill, ln, mknod, mktemp, mount, more, less, netstat, nice, pipe-progress, ps, rm, sync, bc, cal, crc32, dc, diff, du, expand, factor, find, fold, hexdump, install, md5sum, nproc, openvt, pgrep, printf, pscan, pstree, script, seq, sha*sum, shuf, sort, split, strings, tail, tee, uniq, man.
 - Compression/archiving: cpio, gzip, gunzip, lzop, rpm, tar, vi, ar, bunzip2, bzip2, dpkg, lzcat, lzma, unzip, xz.
 - System management: dmesg, setfont, powertop, crond, crontab, lsof, lspci.
 - Networking: ping, udhcpc, tftpd, telnet, sendmail, ntpd, inetd, httpd, ftpd, arping, ftpget, ftpput, nc, traceroute, wget.
 - Editors: ed, sed, awk, vi, hexedit, patch.

End Result: Full list of features (II)

- Kernel: ATA, SATA, PCI, IDE support; MS-DOS filesystems (read/write), mitigations.
- SCSI, E1000 (generic network card interface) support, IPv4 networking, DNS resolution.
- Notably no IPv6 or EXT4 support (they're too bulky).
- Needs at least 70MB of RAM for the decompressor. Not that uncommon among i486/i586 machines.
- Games: 2048, chess; Demos: donut.

Live demo!

- Future ideas:
 - Strip down the BusyBox a bit more to make some space in the initrd for cooler stuff.
 - Add a `.profile` and a cool rice to the floppy disk Linux distribution.
- Programs to bundle:
 - Terminal games.
 - Interpreters and compilers.
 - System rescue tools?
- Package manager:
 - Download a platform toolchain (https://skarnet.org/toolchains/native/i486-linux-musl_i486-14.2.0.tar.xz), compile source packages.

Thanks!

- <https://palaiologos.rocks>
- <https://github.com/kspalaiologos>
- <mailto:kszewczyk@acm.org>
- I hope that you liked this talk. Unfortunately due to real life circumstances and my university obligations - I am currently working on a high-profile thesis about genome compression - I started working on these slides and the code, from scratch, only on Monday 24/03 and was finished by Wednesday 26/03.
- This explains why the decompression performance is not top notch and the boot sector is holey.